

The Free Local AI *Audit* Worksheet

Work through these five sections once and you'll know exactly what to cancel, what to keep, and what to build tonight.

NAME

DATE

ANNUAL AI SPEND \$

1

Audit Your Current AI Spend

- List **every AI tool you pay for monthly** — APIs, Copilot, cloud notebooks, subscriptions.
- For each tool, write the **monthly cost**, the *primary task* it handles, and whether that task touches `sensitive or proprietary data`.
- Multiply each by 12. Sum the column. **That number is your baseline.**

TOOLS / NOTES

2

Install Ollama & Run First Local Inference

- Install Ollama at **ollama.com**, then pull a model and run the job-scoring snippet from the blog.
- Substitute a *real job description* and a short summary of your own background. **Time the response.**
- Note: Is the JSON parseable? What was the latency? Does the output feel usable?

```
ollama pull llama3
response = ollama.chat(model='llama3', messages=[{'role':'user','content': prompt}])
```

LATENCY / OUTPUT

3

Classify Your Tasks: Local vs. Hosted

- Return to your tool list from Section 1. Tag each task **L** (local candidate) or **H** (hosted).
- *Local* = batch, async, runs on a schedule, or touches data that can't leave your machine. *Hosted* = interactive, latency-sensitive, non-sensitive data.
- Calculate: **what percentage of your annual spend lives in the L bucket?**

L / H TAGS

4

Build a Minimal Overnight Automation

- Design a script — *pseudocode is fine* — that runs on a schedule, sends **zero data to external APIs**, and produces a ranked or scored output you'd read in the morning.
- It can be job listings, RSS feeds, a folder of documents — **anything real to you.**
- Write the core Ollama prompt you'd send, and list the `JSON keys` you'd want back.

```
prompt = f"Score this item 1–10. Return JSON: {{score, reasons, flags}}"
```

SCRIPT / PROMPT

5

Measure the Quality Gap on a Real Task

- Pick **one task** you currently send to a paid API. Run the identical prompt against a local `7B` or `13B` Ollama model and against the paid API.
- Score both outputs on **accuracy** and **usefulness** (1–10 each).
- Write two sentences on where the gap is real. Write one sentence on whether *that gap* justifies the cost for this specific task.

SCORES / VERDICT