

Build Your *Autopilot* MLOps Stack

Work through five checks and walk away with a stack sized to your real traffic, not a Netflix fantasy.

NAME

DATE

MONTHLY \$ SAVED

1 Count Your Real Traffic

- Pull **predictions/day** from logs, DB table, or gateway dashboard — last 7 days ÷ 7.
- Write down what you run now: *Kubernetes*, managed endpoint, or cron job.
- Match number to machinery, honestly — 200/day rarely needs a cluster.

PREDICTIONS/DAY

2 Find Your Own Relative-Path Trap

- Grep training + serving code for paths set without `absolute` or explicit env var — `./mlruns` is the classic.
- List every offender found.
- Write one **assert** per path that catches wrong *cwd*.

```
assert ARTIFACT_ROOT.is_absolute()
assert ARTIFACT_ROOT.exists()
```

PATH OFFENDERS FOUND

3 Rewrite Your Cron Line

- Open crontab / systemd timer for your training job.
- Check two rules: full path to venv `python`, `MAILTO` or webhook on failure.
- Rewrite the missing piece, test from `$HOME` to confirm it still resolves.

```
0 3 * * * MAILTO=you@example.com /home/you/proj/.venv/bin/python /home/you/proj/train.py
```

CRON LINE FIXED?

4 Pick Your Tool Swaps

- List your stack: **orchestrator**, **tracking**, **serving**, **monitoring**.
- For each paid or over-provisioned tool, name the free swap — **Prefect**, self-hosted **MLflow**, **FastAPI+Uvicorn**, **Evidently**.
- Estimate the monthly *dollar* difference per swap.

\$ SAVED PER SWAP

5 Decide What You're Allowed to Skip

- Name one check you're skipping — a *Parquet log* instead of full drift monitoring, say.
- Write the exact condition that forces you to build it **for real**.
- Weigh cost of skipping against cost of building — write both down.

TRIGGER CONDITION